

Aaditya Pandagle

+91 7972567926 | adi.pandagle@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Full Stack and AI/ML Engineer with 2+ years building production SaaS platforms, Retrieval-Augmented Generation (RAG) systems, and enterprise automation pipelines. Proven track record of delivering measurable impact across B2B and large-scale engineering environments.

SKILLS

Languages: Python, TypeScript, SQL
Full Stack: React, Next.js, Node.js, TailwindCSS, Git, FastAPI, Django, Prisma, Redis
AI/ML: LangChain, LangGraph, RAG, Vector Search, Hugging Face, Ollama, TensorFlow, PyTorch, Generative AI
Cloud/Databases: AWS, Docker, PostgreSQL, MongoDB, Qdrant, MLflow, Jenkins
Practices: REST API, CI/CD, Microservices, Agile

WORK EXPERIENCE

Project Engineer <i>Wipro (ABB)</i>	Nov 2024 – Present <i>Pune, India</i>
- Engineered Python/Bash automation scripts to streamline repetitive workflows, measurably reducing manual monitoring overhead and improving deployment infrastructure efficiency. - Led root-cause analysis on production failures by parsing system logs and isolating runtime bottlenecks, driving significant MTTR reduction for critical L3/L4 engineering escalations. - Optimised data workflows and API communication within complex enterprise architectures, ensuring seamless cross-system integration and improving runtime stability across critical services. - Partnered with senior developers to build reproducible sandbox environments for production defect isolation, accelerating delivery cycles and strengthening enterprise system reliability.	
Data Science Intern <i>NSArrows Innovations</i>	June 2024 – Sep 2024 <i>Pune, India</i>
- Built a FastAPI + MongoDB service to extract and structure data from 500+ unstructured PDF documents, automating manual ingestion workflows and reducing data processing time by 40%. - Developed a Python pipeline for automated video transcription and multilingual voice synthesis across 3+ languages, cutting internal content preparation time by 65%. - Designed a semantic search engine using Qdrant and MongoDB vector embeddings for internal knowledge retrieval, improving search relevance and reducing time-to-find across content libraries.	

KEY PROJECTS

HealthCare Foresights healthcareforesights.com	Nov 2025 – Present
- Built a B2B SaaS platform using Next.js 15, React 19, and TypeScript to stream healthcare market research reports to 50+ enterprise clients, with secure gated previews and a Lighthouse score of 95+. - Architected a Go (Fiber v2) REST API with clean architecture, PostgreSQL (GORM) for persistence, and Redis caching, achieving sub-100ms endpoint latency under concurrent load. - Implemented PayPal REST API payment workflows with server-side validation, Redis token caching, and async email notifications, processing \$6K+ in monthly transactions. - Built an admin dashboard with TailwindCSS and ShadCN consolidating report CMS, order management, and real-time sales analytics, reducing admin overhead by 50%.	
GitTalk github.com/iamblesalt/GitTalk	Oct 2025 – Present
- Built a local-first AI codebase assistant utilising React 19, React Router v7, and FastAPI, enabling intelligent conversational Q&A over repositories up to 500MB with sub-second response latency. - Parsed multi-language repos via tree-sitter AST into structural chunks, implementing a hybrid search engine (BM25 + vector) with LanceDB, achieving 90%+ retrieval accuracy on code queries. - Integrated Ollama (qwen2.5-coder) for fully offline reasoning with a SHA-256 incremental re-indexing pipeline, accelerating subsequent scans by 70% and eliminating redundant processing.	
Agentic RAG Workspace github.com/iamblesalt/musical-octo-spoon	March 2026 – Present
- Architected a local-first agentic RAG system with zero data egress, organising personal documents into isolated workspaces across 7+ formats (PDF, DOCX, EPUB, MOBI, FB2, CBZ, XPS). - Designed a tool-calling agent loop with adaptive search-budgeting, query rephrase-and-retry, and cross-encoder reranking, layering hybrid dense + sparse (BM25) retrieval for low-latency search across 50K+ pages. - Implemented an automated faithfulness/citation-accuracy eval loop that scores generated answers against retrieved passages, catching ungrounded claims and grounding every answer in page-level source citations. - Delivered a real-time Next.js chat UI with SSE live tool-activity streaming and local on-disk observability logging, giving full traceability into every retrieval and tool call.	

EDUCATION

Post-Graduation Program in Data Science and Analytics <i>Imarticus Learning (TIER-2)</i>	2023 – 2024 <i>Pune, India</i>
Bachelor of Business Administration (Computer Application) <i>Modern College of Commerce & Computer Studies, Savitribai Phule Pune University</i>	2020 – 2023 <i>CGPA: 8.61</i>

CERTIFICATIONS

Udemy: <u>Python for Data Science & ML Bootcamp</u> • <u>Python Data Structures and Algorithms</u> • <u>AWS Solutions Architect</u> • <u>AWS Generative AI Developer Professional</u>
Kaggle: <u>Intro to Machine Learning</u> • <u>Intermediate Machine Learning</u>
HackerRank: <u>SQL Basic</u> • <u>SQL Intermediate</u>